

Translating Polish Sermons Cheaply: LLM APIs vs Dedicated MT (2026 Cost & Quality)

TL;DR

For a single weekly Polish→English sermon (~1 hour), the translation step is effectively free-to-trivial: a cheap LLM (Gemini 2.5 Flash-Lite, GPT-5 nano, Qwen Flash, DeepSeek, or Mistral Small) costs roughly **\$0.006–\$0.04 per hour** even with context-carrying, i.e. **under ~\$5/year**; a dedicated MT service costs **\$0.55–\$2.70/hour** on paid tiers but is **\$0/year** if you stay inside free tiers (DeepL 500k, Google 500k, Azure 2M chars/month all cover a weekly service).

On quality, Polish→English is the *easy* direction, and both approaches are strong: DeepL and modern LLMs are essentially tied at the top for European/Slavic languages (nearest rigorous proxy, the Slovak skLEP benchmark, rated DeepL 3.70/3.73 vs GPT-4o 3.62/3.73 vs Google 3.57/3.67 on fluency/adequacy), while independent 2025 testing (Intento) found LLMs were the top performer in **89% of language pairs**.

Bottom line: choose a cheap LLM with a strong system prompt. Its unique ability to enforce a church glossary, fixed Bible-reference formatting, tone/register, and cross-segment consistency matters far more than the near-zero cost difference — with DeepL's free tier as an excellent quality fallback and a human reviewer for doctrinally sensitive passages.

Key Findings

The translation step is nearly free at this volume. A one-hour sermon is only ~8,500 words / ~55,000 characters / ~15,000 Polish input tokens. Cheap LLMs price input/output at cents per million tokens, so per-hour cost is fractions of a cent to a few cents. Dedicated MT bills per source character at \$10–\$50 per million characters — a much higher *effective* rate — yet still only ~\$0.55–\$2.70/hour, and \$0 inside free tiers.

LLMs are ~20–100× cheaper per hour than dedicated MT on raw usage cost, because per-character MT pricing (~\$20/M chars) works out to a far higher effective per-token rate than cheap-LLM token pricing (~\$0.10/M tokens). The trade-off runs the other way on speed: per Intento's 2025 analysis, LLMs are "priced by token rather than character and are 10 to 100 times slower than NMT engines for the same content volume." Context-carrying multiplies LLM input several-fold but the absolute cost stays trivial.

Quality for Polish→English is strong across the board. LLMs handle the *comprehension→English generation* direction well; the morphological difficulty of Polish mainly hurts the *reverse* (English→Polish) direction. DeepL remains the strongest dedicated engine for European languages; GPT-4.1/Gemini 2.5 Pro/Claude now match or exceed it in independent 2025 testing.

Steerability is the LLM's decisive advantage for scriptural content: a free-form system prompt can enforce a glossary of church terms, proper-noun spellings, a consistent Bible-reference format, register/tone, and cross-segment context — none of which dedicated MT services fully replicate.

Details

Volume model (assumptions stated)

Sermon delivery ~130–150 wpm (the standard preaching band; experienced preachers naturally land around 140 wpm, ~4,200 words per 30 minutes); continuous speech ≈ **8,000–9,000 words/hour** — I use **8,500 words/hour**.

English output ≈ same word count (~8,500 words).

Characters (for per-character MT, billed on **source/Polish** text): ~6–7 chars/word → **~55,000 Polish characters/hour**.

Tokens: Polish tokenizes inefficiently on Western LLM tokenizers (Bielik research measured ~3.2 tokens/word on Llama-2/Mistral tokenizers vs ~1.6 with a Polish-optimized tokenizer; modern o200k/Gemini tokenizers land in between). I use **~1.8 tokens/Polish word** → **~15,000 input tokens** of source text and **~12,000 output tokens** (English, ~1.33 tok/word).

Scenario A (naive): each segment translated once with a small fixed prompt → I model **~20,000 input + 12,000 output tokens/hour** (raw text plus light per-call system-prompt

overhead).

Scenario B (context-carrying): previous source+translation segments re-sent every call for consistency, $\sim 5\times$ input inflation (within the user's $3\text{--}10\times$ range) $\rightarrow \sim \mathbf{100,000}$ input + $\mathbf{12,000}$ output tokens/hour. Prompt caching (10% of the input rate on OpenAI/Anthropic/Google/DeepSeek) can cut most of the re-sent context cost.

1) Cost per hour — general-purpose LLM APIs (cheap tiers)

Rates below are per **million tokens** (input / output), current 2026:

Model	Input / Output (\$/M)	Naive: ~\$/hour	Context-carry: ~\$/hour
OpenAI GPT-5 nano	\$0.05 / \$0.40	\$0.006	\$0.010
OpenAI GPT-4o-mini	\$0.15 / \$0.60	\$0.010	\$0.022
OpenAI GPT-5 mini (sunsets Dec 11, 2026)	\$0.25 / \$2.00	\$0.029	\$0.049
OpenAI GPT-5.6 Luna	\$0.20 / \$1.20	\$0.018	\$0.034
Google Gemini 2.5 Flash-Lite	\$0.10 / \$0.40	\$0.007	\$0.015
Google Gemini 2.5 Flash	\$0.30 / \$2.50	\$0.036	\$0.060
Anthropic Claude Haiku 4.5	\$1.00 / \$5.00	\$0.080	\$0.160
DeepSeek V4 Flash / chat	\$0.14 / \$0.28	\$0.006	\$0.017
DeepSeek V3.2	\$0.28 / \$0.42	\$0.011	\$0.033
Alibaba Qwen Flash (≤256k, via Model Studio/DashScope)	\$0.05 / \$0.40	\$0.006	\$0.010
Mistral Small (4)	\$0.15 / \$0.60	\$0.010	\$0.021
Groq Llama 3.3 70B	\$0.59 / \$0.79	\$0.021	\$0.068
Together Llama 3.3 70B	~\$0.88 flat	\$0.028	\$0.099
Even the priciest "cheap" option (Claude Haiku 4.5) is ~\$0.08–0.16/hour. The genuinely cheap tier (Gemini Flash-Lite, GPT-5 nano, Qwen Flash, DeepSeek, Mistral Small) is well under \$0.02/hour naive. Note that per Intento/MachineTranslation.com's 2026 comparison, DeepSeek is "the most cost-efficient LLM API option" — at roughly \$0.14–0.28/M tokens it is essentially tied with Gemini Flash-Lite and Qwen Flash for cheapest viable choice. (Qwen's			

older "Turbo" tier is being retired in favor of Flash; Mistral "Small 4" is the current small tier at \$0.15/\$0.60, though some trackers still list a legacy Small 3 at \$0.10/\$0.30.)

2) Cost per hour — dedicated MT APIs

Priced per **million source characters** (~55,000 chars/hour):

Service	Rate	~\$/hour	Free tier
Google Cloud Translation — NMT (Basic v2 / Advanced v3)	\$20/M chars (input only)	~\$1.10	500k chars/month (permanent)
Google — Translation LLM	\$10/M in + \$10/M out	~\$1.10	Not covered by free tier
Google — Adaptive Translation	\$25/M in + \$25/M out	~\$2.70	Not covered by free tier
Google — AutoML custom NMT	\$80/M chars	~\$4.40	Separate 500k/month allotment
DeepL API (current Growth / legacy Pro)	~\$25–27.50/M chars (+ small base fee)	~\$1.38	Free tier: 500k chars/month
Microsoft Azure AI Translator (S1)	\$10/M chars	~\$0.55	2M chars/month (F0)
Amazon Translate	\$15/M chars	~\$0.83	2M chars/month (first 12 months)

Azure is the cheapest paid dedicated MT (~\$0.55/hr); Google Adaptive/AutoML are the most expensive. Google's own examples confirm the Translation LLM tier "works out to ~\$20/1M — same effective price as NMT, but with better quality and higher latency." **See the DeepL pricing caveat below** — sources conflict and DeepL restructured its API plans in mid-2026.

3) Quality for Polish→English (religious/formal text)

Direction matters. Polish→English is the easier direction. FLORES/WMT research systems

score Pl→En ~28–34 BLEU vs En→Pl ~19–27 BLEU — a consistent 6–8 BLEU gap (e.g., WMT20 Volctrans: Pl→En 34.4 vs En→Pl 26.1). Because your pipeline is Pl→En, you are on the favorable side.

DeepL vs Google vs LLMs. DeepL has historically been the top engine on ~65% of language pairs (Intento), concentrated in European languages — the basis for the long-standing "use DeepL for European languages" recommendation. The closest rigorous Slavic head-to-head is the Slovak skLEP benchmark (Suppa et al., arXiv:2506.21508 / ACL Findings 2025), which rated systems on 4-point fluency/adequacy scales: **DeepL 3.70/3.73, GPT-4o 3.62/3.73, Google Translate 3.57/3.67**, MADLAD-400-3B 3.48/3.53, NLLB-3.3B 3.40/3.54 — statistically a near-tie among the top three.

LLMs now lead most pairs. Intento's *State of Translation Automation 2025* (46 engines, 11 pairs, human LQA) found "LLMs represent 89% of top performers across language pairs," with "GPT-4.1 leads with 7 'best' performances across 11 language pairs, followed by Gemini 2.5 Pro (5) and Lara (4)"; DeepL's newer LLM-based ("next-gen") engine was the best of the dedicated real-time systems. The WMT24 General MT findings (Kocmi et al., statmt.org) ranked **Claude 3.5 Sonnet first, winning 9 of 11 language pairs**, ahead of GPT-4 and every online provider. There is no peer-reviewed table isolating Polish for these commercial systems, but the European/Slavic evidence points the same way.

Morphology caveat. LLMs measurably degrade on Polish's rich inflection (studies show +14–23% accuracy gains for Polish inflection when given explicit linguistic rules), but this primarily affects En→Pl *generation*, not the Pl→En direction you need.

Hallucination caveat. LLMs can subtly embellish, paraphrase, or fabricate. A GDELT study cited by MachineTranslation.com found "Gemini 2.5 Pro produced different translations of the same Russian passage on consecutive runs — including hallucinated organisation names." For scripture and doctrine this is a real risk — keep a reviewer and instruct the model to translate faithfully without adding content.

4) Customization / steerability

LLM (free-form prompt) — unique capabilities: enforce a church glossary and proper-noun spellings; convert Polish Bible-reference style (e.g., "J 3,16") to English style ("John 3:16") consistently; preserve poetic/liturgical register; carry the previous segments as context for coherence; give explicit handling rules for ambiguous words — all with **zero setup**, just

editable text. No dedicated MT service can reformat references, preserve register on command, or reason about ambiguity.

DeepL: glossaries via the REST API (one-to-one term pairs), plus formal/informal tone control; Polish is supported for glossaries and as source/target. Predictable but limited — no free-form instructions.

Google Cloud: glossaries + AutoML custom models + Adaptive Translation (few-shot steering from your own examples). More setup, added cost (\$25+\$25/M for Adaptive, \$80/M for AutoML).

Azure Custom Translator: dynamic dictionary (safe mainly for proper/compound nouns, requires English on one side), [microsoft](#)

↗
phrase dictionary, or a full custom model (needs ~10,000 parallel sentences; Microsoft cites 5–10+ BLEU gains with enough data). Heaviest setup.

Amazon Translate: custom terminology via uploaded CSV.

5) Weekly / annual budget

Cheap LLM, naive: $\sim \$0.01/\text{hr} \times 52 \approx$ **\$0.50/year**. Context-carrying: $\sim \$0.03/\text{hr} \times 52 \approx$ **\$1.50/year**. A 1.5-hour full service scales these by 1.5 → still only a few dollars/year.

Dedicated MT paid: DeepL $\sim \$1.38/\text{hr} \times 52 \approx$ **\$72/year**; Google NMT $\sim \$1.10 \times 52 \approx$ **\$57/year**; Azure $\sim \$0.55 \times 52 \approx$ **\$29/year** — but see free tiers.

Free tiers cover a weekly service outright. One 1-hour service $\approx$ 55k chars; ~ 4.3 services/month $\approx$ $\sim 235\text{k}$ chars — inside DeepL's and Google's 500k/month free allotments (per DeepL's Help Center, "With DeepL API Free, you can translate up to 500,000 characters per month for free"), and far inside Azure's 2M/month. Even a weekly 1.5-hour service ($\sim 355\text{k}$ chars/month) fits DeepL/Google free tiers. So a single weekly service can run at **\$0** on a dedicated MT free tier.

6) Latency / real-time suitability

Fastest for streaming: dedicated NMT (Google NMT, Azure, Amazon) is millisecond-latency and purpose-built for real-time. Among LLMs, Groq (Llama 3.3 70B at $\sim 276\text{--}394$ tok/s, sub-100ms first token) [Eesel AI](#)

↗
[Spheron](#)

and Gemini Flash-Lite / GPT nano are fast enough for near-real-time segment translation.

Slower: Google's Translation LLM is $\sim 3\times$ slower than its NMT; [ChatsControl](#)

context-carrying LLM calls add latency because more input must be processed each turn; and LLMs are generally an order of magnitude or more slower than NMT engines per Intento.

Rate limits: free LLM tiers throttle (e.g., Groq's free tier is 30 requests/min, 1,000/day). A sermon chunked into many short segments can hit RPM caps — batch a few seconds of text per call. Qwen-Flash offers very high throughput (multi-million TPM) if you need many small calls.

Recommendations

Primary: a cheap LLM with a curated system prompt. Start with **Gemini 2.5 Flash-Lite** or **GPT-5 nano** (both $\sim \$0.006\text{--}0.015/\text{hour}$), or **Qwen Flash / DeepSeek V4 Flash / Mistral Small** if you prefer. Put your church glossary, proper-noun spellings, Bible-reference format, register instructions, and a "translate faithfully, add nothing" rule in the system prompt. Carry the last 2–4 segments as context for coherence; enable prompt caching to keep the re-sent context near-free. Expected annual cost: **\$1–\$5**.

Quality fallback / A-B check: DeepL free tier (500k chars/month). It's free for a weekly service and is the strongest dedicated engine for Polish. Use it to sanity-check the LLM's output and as a hot backup if the LLM endpoint is down.

If near-real-time latency is critical: use Gemini Flash-Lite or Groq for the LLM path; where steerability matters less, Google NMT/Azure give the lowest latency.

Always keep a human reviewer for doctrinally sensitive passages, given LLM hallucination risk.

Benchmarks/thresholds that would change this advice:

If you switch to **English→Polish** (the hard direction): lean on DeepL and/or add human post-editing; LLM morphology errors rise.

If **hallucination/fidelity** proves unacceptable in review: move to dedicated NMT (Azure/Google) or mandate human post-edit.

If volume grows to **many services/day** or long archives: you'll exit free tiers — the cheap LLM's per-hour cost still wins by a wide margin; watch RPM/TPM rate limits and batch accordingly.

If you need **guaranteed terminology consistency** with zero prompt-engineering: a DeepL/Azure glossary is more predictable than prompt instructions.

Caveats

DeepL API pricing is in flux and sources conflict. Multiple 2026 sources cite DeepL API Pro at "\$5.49/month + \$25 per 1M characters"; [ChatsControl](#)

↗ others cite "\$5.49 per 1M characters"; [NeuraPlusAI](#)

↗ and it is confirmed (DeepL Help Center; Langbly, Jan 30 2026) that in **July 2026 DeepL retired API Free and API Pro** for new purchases, steering new customers to **Developer** (1M chars total, one-time free) and **Growth** (\$26/month including ~12M chars/year, then ~\$27.50/M overage, up to a 50M chars/month cap). The 500k/month free allowance still applies to existing/free API keys. Verify the live plan before committing; my ~\$1.38/hour figure uses the ~\$25/M rate.

All per-hour figures are estimates built on the stated 8,500-word / 55,000-character / token assumptions. Real sermons vary in pace, pauses, and length; a slower or shorter sermon costs proportionally less.

Polish token counts are approximate — Polish is tokenized less efficiently than English on most Western LLM tokenizers, so input-token cost per word is higher than the English rule of thumb; [NeuralMindMastery](#)

↗ measure on your own model and text.

Pricing changes fast. Model names and rates (especially the OpenAI GPT-5.x, Gemini 3.x, and Qwen 3.x tiers) shift monthly, and Groq moved its headline Llama models to "contact sales" in late August 2026; several cited figures come from pricing-tracker sites rather than the vendor's own page. Re-check vendor pricing pages before finalizing a budget.

No public benchmark isolates Polish→English religious/scriptural text. Quality claims for that specific register are inferred from general Polish/Slavic MT quality plus LLM steerability, not a direct scriptural benchmark.

Whisper transcription quality is the upstream bottleneck. Translation quality is capped by transcription accuracy; Polish ASR errors (proper names, theological terms) will propagate — a glossary in the LLM prompt can partially correct these, another point in favor of the steerable LLM approach.